Venture Name: MedBrief
Venture Website
Venture Short Description (max 250 char)
MedBrief is an AI platform for personal injury and insurance lawyers that reliably turns complex medical records into verified, litigation-ready evidence.

Venture Long Description (150 words)
MedBrief is an AI platform for personal injury and insurance lawyers that turns complex medical records into verified, litigation-ready evidence. Legal, insurance, and medical organizations still rely on slow, manual processes to extract meaning from handwritten, faxed, and scanned medical files, while existing OCR and generic AI tools fail to produce reliable results.

MedBrief uses specialty-trained medical AI to extract, verify, and structure clinical information across large volumes of unstructured records. Every data point is cross-checked, assigned a confidence score, and linked back to its source, creating an auditable record suitable for discovery, expert review, and trial. The verified data is used to generate medical summaries and case-specific legal arguments enriched with relevant legal standards.

Lawyers and medical reviewers can interact with both the underlying records and the generated reports through a conversational interface, making medical-legal document review faster, more consistent, and more defensible.

PROBLEM & SOLUTION

Describe the problem you are trying to solve. What is the need? (300 words)
Personal injury litigation relies heavily on medical evidence, yet translating clinical records into a litigation-ready argument is time-consuming, expensive, and inefficient. Personal injury firms routinely receive hundreds or thousands of pages of medical records per case. These records are filled with handwritten notes, images, charts, upside-down faxed pages, and scanned reports produced by multiple providers. These records are unstructured, duplicative, and often unsearchable due to limitations in image- and handwritten-text recognition.

Lawyers are not trained to interpret raw clinical data, while medical professionals are not trained to identify which facts are material to litigation. As a result, the current system relies on physicians manually summarizing medical records into narrative reports in layterms for attorneys. This process is time-intensive, inconsistent in quality, and misaligned with the needs of litigation and claims processing. Reports often take weeks or months to produce, delaying case evaluation, settlement negotiations, and insurance payouts.

General-purpose AI tools have not solved this problem. Large language models and conventional document-processing systems struggle with handwritten notes and domain-specific clinical terminology. More importantly, large LLMs are known to hallucinate, so they lack the precision and source traceability required in regulated legal and medical settings. Lawyers cannot rely on outputs that cannot be verified against the underlying medical record, and doctors cannot risk generating summaries that misstate clinical facts.

The need is clear: personal injury firms and the physicians they work with require a faster, more accurate, and verifiable way to extract and translate medical evidence into legally actionable information. Without a system that bridges medical data and legal reasoning, inefficiency and preventable claim failures will continue to undermine fair outcomes for injured individuals and delay access to compensation and care that the personal injury system is meant to provide.

Describe your proposed solution. How would this solve the problem outlined above? (300 words)

MedBrief addresses the inefficiencies that slow and fragment personal injury litigation. It is a medical-legal intelligence platform that converts unstructured clinical records into legally actionable evidence, compressing what traditionally requires tens of hours of physician time and legal review into minutes.

First, MedBrief performs medical-grade document data extraction. The platform takes in heterogeneous medical records, including handwritten notes, scanned charts, and diagnostic reports, and converts them into structured, machine-readable data. Unlike general-purpose OCR or large language models, MedBrief is optimized for clinical documentation, preserving medical terminology, temporal context, and relationships across records. This step eliminates the need for lawyers or physicians to manually sift through hundreds of pages of unstructured material.

Second, MedBrief generates auditable, verifiable medical reports that are both clinically accurate and legally relevant. Extracted information is organized into structured summaries that include personal injury case-relevant facts such as injury chronology and treatment history. Every piece of reported data is traceable to its exact location in the source record, allowing physicians to review, verify, and approve outputs with confidence. Any portions of the record where MedBrief has lower confidence in extraction accuracy are explicitly flagged for review rather than skipped or inferred. This ensures medical precision while aligning reports with the evidentiary needs of litigation and claims evaluation.

Third, MedBrief provides legal professionals with an informed, case-specific chatbot that generates evidence-based legal arguments for personal injury litigation.
Together, these steps reduce report turnaround times, lower costs for firms and physicians, and ensure that valid claims are evaluated on their medical merits, supporting fairer and faster outcomes for injured claimants.

INNOVATION & IMPACT (300 words each)

How are users currently solving this problem?

Personal injury firms currently address the challenge of medical evidence review through a fragmented mix of manual effort, outsourcing, and point solutions that each cover only a small portion of the overall workflow. Lawyers and paralegals manually review hundreds to thousands of pages of unstructured medical records (hospital records, physiotherapy reports, diagnostic imaging reports, fracture clinic reports, etc.) to extract the facts needed for case development. This process is slow, costly, and highly dependent on individual expertise.

At the same time, firms and insurance companies frequently refer claimants to physicians they work with for evaluation. These doctors prepare summary reports, such as chronic pain assessments, to support insurance decisions and litigation. Producing these reports takes time. They are often written without full alignment to legal evidentiary needs.

To manage this burden, some firms outsource medical review to legal nurse consultants or third-party vendors. This can improve quality. It also introduces new delays and high costs. Many smaller firms cannot rely on these services consistently and instead absorb the work internally.

Based on firsthand experience at a law firm that deployed an LLM built on OpenAI's ChatGPT, these tools have proven limited in practice, especially for in-depth analysis and data synthesis. They struggle to reference long documents and lack the specificity required for reliable analysis. Similar limitations were observed by colleagues across the firm. These AI tools provide general guidance on performing legal tasks, but they do not perform the underlying work, limiting their usefulness for experienced attorneys.

As a result, firms stitch together people, vendors, and partial tools to move cases forward. The effectiveness of this approach depends heavily on staffing, budgets, and turnaround times, leaving many valid claims underdeveloped or abandoned due to the cost and complexity of medical evidence review.

How would your solution be an improvement from the status quo?

MedBrief improves on the status quo by addressing the gap between physician-generated reports and the underlying medical records that ultimately determine case strength. Today, lawyers rely heavily on summary reports prepared by physicians to evaluate claims and communicate with insurers. These reports are valuable, but they are necessarily selective. They cannot capture every medically relevant detail contained in the full record, nor are they designed to anticipate every legal issue that may arise as a case develops.

As a result, lawyers often need to revisit the source medical files to answer new questions, test theories, or respond to challenges from opposing counsel or insurers. Doing so requires manually navigating large volumes of unstructured records that are difficult to comprehend without tools designed to support legal analysis at the source document level.

MedBrief changes this by operating directly on the underlying medical records while preserving the value of physician review. Instead of replacing physician reports, MedBrief augments them. It extracts and structures medical data from source files, allows physicians to verify its accuracy, and enables legal analysis that remains directly linked to the original documentation. Lawyers can analyze medical evidence based on the full record, not just a summary, and trace every conclusion to its source.

This approach allows firms to evaluate claims more thoroughly and earlier in the process. Legal arguments and issue summaries are generated based on verified medical facts drawn from the complete record, rather than relying solely on static reports. The result is faster case development, lower costs, and greater confidence that litigation decisions are grounded in the strongest available medical evidence.

By enabling source-based analysis without requiring manual review, MedBrief improves both efficiency and accuracy, ultimately supporting better outcomes for lawyers, physicians, and injured claimants.

If successful, what would be the scale and depth of your venture's impact/effect on the market?

If successful, MedBrief would reshape how medical evidence flows through the personal injury system. At the firm level, lawyers would spend significantly less time on administrative medical record review, which today can account for 20-40% of case preparation time, and more time on advocacy and strategy. Physicians would reclaim hours currently spent compiling narrative reports, a task that can take 10-30 hours per complex case, while maintaining confidence in the accuracy and completeness of the medical record. Claims could be evaluated earlier, reducing delays that often stretch case timelines by weeks or months.

At the market level, MedBrief would lower the economic threshold for pursuing valid claims. Firms could take on cases that would otherwise be declined due to medical review costs exceeding expected recoveries, thereby increasing access to legal representation for injured individuals. Faster, evidence-based evaluations would also support earlier settlements and more efficient dispute resolution.

At a broader societal level, improving the translation of medical evidence into legal outcomes strengthens trust in the civil justice system. When claims are resolved based on verified facts and not hindered by procedural barriers, injured individuals are more likely to receive fair and timely compensation. Over time, this reduces inequities driven by resource constraints rather than merit.

Beyond personal injury, the same infrastructure can extend to adjacent markets, including insurance disputes, workers' compensation, disability claims, and mass tort litigation. By serving as a foundational medical-legal intelligence layer, MedBrief has the potential to become core infrastructure for evidence-based decision-making at the intersection of medicine and law.

MARKET (300 words each)

Describe the customer segment you are planning to target first. Why is this the right market entry point?

MedBrief's initial target customer segment is legal buyers, specifically personal injury law firms, medical-legal practices, and litigation support teams that routinely handle large volumes of medical records. These users sit at the intersection of high evidentiary standards, time-sensitive decision-making, and complex unstructured medical data, making them an ideal entry point for the platform.

Legal professionals face immediate and recurring pain points in medical evidence review. Case outcomes depend heavily on accurate medical interpretation, yet the underlying records are often handwritten, faxed, fragmented across providers, and poorly structured. Existing tools either stop at basic document management or provide generic AI assistance that lacks medical specificity, auditability, and legal defensibility. As a result, lawyers rely on slow, manual review processes and static physician reports that are costly to produce and difficult to interrogate as cases evolve.

Legal buyers are the right initial market for three reasons. First, the value proposition is immediate and measurable. Time saved in medical review directly translates into lower costs, earlier case evaluation, and stronger litigation positioning. Second, legal workflows demand traceability, and source-level audit trails which are all requirements that align closely with MedBrief's core technical differentiation. Building to these standards early creates a defensible foundation for other regulated markets. Third, legal adoption is relationship-driven and bottom-up. Individual lawyers and small firms can adopt the platform quickly, confirm value on real cases, and expand usage organically within firms.

Starting with legal buyers also creates a natural expansion path. The same verified medical intelligence that supports litigation can be extended to insurance claims, disability adjudication, and medical enterprise workflows, where similar data challenges exist at larger scale. By first proving value in a high-stakes, evidence-driven legal environment, MedBrief establishes credibility and product maturity before moving into broader institutional markets.

Who are the key competitors? How will your product or service be an improvement over what competitors offer?

Key competitors fall into three broad categories:

1) General-purpose AI platforms: Tools such as OpenAI, Google Gemini, and Anthropic Claude are strong at text generation and summarization. However, they are not designed for medical record extraction or legal use. They struggle with handwritten notes and scanned documents and are known to hallucinate. They also lack verification, confidence scoring, and reliable source traceability. As a result, their outputs are difficult to trust in regulated legal and medical settings. Google has also released MedGemma, a set of open medical AI models. While MedGemma improves general medical understanding, our testing on real medical records showed that it remains insufficiently precise for review. It does not provide specialist-level reasoning, confidence indicators, or end-to-end auditability. In practice, it functions as a medical assistant rather than a defensible medical-legal system.

2) Document AI and OCR platforms: Tools such as Landing AI focus on document extraction. They perform well on clean, structured inputs. However, they break down when applied to real-world medical records that include handwriting, poor-quality faxes, or mixed layouts (pictures and text). With poor data extraction, features such as a chatbot for discussing and reviewing the content are not helpful.

3) Legal drafting and research tools: Research platforms such as Westlaw and LexisNexis answer legal questions, but they do not analyze case-specific medical records. There are also legal AI tools that help polish briefs or demand letters, but they operate downstream. They assume attorney notes already exist.

MedBrief is built to bridge medicine and law. It extracts medical data using specialist-trained models, verifies results using confidence scoring, and links each output to its source. It then generates case-specific legal summaries and arguments grounded in verified medical facts. Instead of stitching together multiple tools, it is a single system designed for accuracy, and trust.

How have you engaged with potential users/customers? How have you validated customer demand?

Yes. MedBrief has engaged directly with both medical professionals and personal injury lawyers throughout product development. The medical and legal sides of the product are being developed in parallel.

On the medical side, we have worked with practicing physicians who operate both private practices and hospital-based clinics. Physicians provided real patient files for testing and continue to send batches of medical records for processing. MedBrief does not yet have an external-facing physician portal. Physicians upload records through cloud storage. We process the documents using our extraction system and return structured summaries for review.

Over the past three months, physicians have repeatedly requested additional batches for processing. They have expressed enthusiasm for the extraction's reliability and the usefulness of the summaries. This repeated usage has helped validate demand for accurate medical document extraction across multiple specialties. It has also confirmed that physicians value shifting time from manual compilation to verification and review.

One of MedBrief's co-founders is a practicing physician with an established surgical practice and hospital affiliations. Through this network, additional doctors have tested the product and provided ongoing feedback. This has allowed the team to iterate quickly and validate demand beyond a single practice.

On the legal side, the team has engaged several attorneys through an AI-based client intake and case assessment tool developed by a co-founder. Lawyers have used the system to conduct claimant interviews and generate issue summaries. The tool references legal documents and evidence to support the arguments it generates. Feedback has been mixed. Some firms use the tool daily and find it valuable. Others noted that intake work is often handled by paralegals. These conversations have been critical in refining product focus and identifying where MedBrief provides the most value. Follow-up meetings are scheduled with firm partners and attorneys to guide further iteration.

How have you validated the viability of your venture idea as a business?

MedBrief's pricing is grounded in the time and cost lawyers and physicians already incur in producing and reviewing medical reports. Our discussions focused on how long this work takes today, what it costs, and whether users would be willing to pay to reduce that burden.

On the medical side, we have spoken with physicians across clinics and hospitals who routinely prepare assessment reports for insurance and litigation. These reports often require consolidating dozens of medical documents and can take many hours to complete. In one example, a physician charges approximately $1,400 to produce a single assessment report that synthesizes records from multiple providers (e.g., in one case, 30 reports, including hospital records, physiotherapy reports, fracture clinic reports, diagnostic imaging results, and blood work). The assessment was conducted in mid-October, and the finalized report was delivered in early December, reflecting weeks of delay driven largely by document review and compilation. Multiple physicians indicated they would be willing to pay on the order of $100 per report to save several hours of manual work, provided accuracy and source traceability were maintained. These conversations align with MedBrief's core technical advantage: automated extraction and structuring of complex medical records with built-in verification and full traceability to source documents.

On the legal side, we examined firm-level economics with personal injury practitioners. A mid-size firm may employ roughly 20 lawyers, each managing more than 100 active cases per year, with each case having over 100 pages of medical records. Costs associated with medical evidence review build up quickly at this scale. Lawyers emphasized that faster access to structured, reliable medical information would increase productivity and allow earlier, more confident decisions on cases that proceed to litigation. This business is viable through either an annual subscription fee for a firm or a per-case fee.

Can you describe other possible markets or customer segments? If yes, describe.

While MedBrief is initially focused on personal injury firms and the physicians they work with, the underlying platform is applicable to a broader set of markets where complex medical records must be reviewed, interpreted, and translated into high-stakes decisions. These adjacent markets share the same core challenges: unstructured medical data, time pressure, and the need for accuracy and auditability.

One natural extension is the insurance sector, including accident and health, disability, and life insurance carriers. These organizations review large volumes of medical documentation to assess claims, determine coverage, and resolve disputes. Medical record review is a major operational bottleneck, often driving long claim cycle times and inconsistent outcomes. MedBrief's ability to extract, verify, and trace medical facts back to source documents makes it well-suited for claims adjudication, internal audits, and dispute resolution.

Another adjacent market is workers' compensation and occupational injury claims. These cases rely heavily on medical timelines, causation analysis, and functional assessments. Delays and ambiguity in medical evidence can stall claims and increase litigation risk. MedBrief can support faster, more consistent evaluation by structuring medical records around legally relevant issues while preserving physician oversight.

MedBrief also has applications in medical-legal services beyond personal injury, including mass tort litigation, disability appeals, and medical malpractice defense. In these contexts, teams must analyze extensive medical histories across multiple providers and time periods. Source-level traceability and auditability are particularly important, making MedBrief's approach well aligned with evidentiary requirements.

Finally, MedBrief can serve medical enterprises themselves, including hospitals, clinics, and health systems, by reducing the administrative burden associated with record reconciliation, external requests, and reporting. By converting fragmented medical documentation into structured, verifiable data, MedBrief can support compliance, utilization review, and downstream analytics.

PRODUCT DEVELOPMENT (300 words each)

Outline your product development plan.

MedBrief's core AI technology has already been built and validated. The following core capabilities are already implemented and tested on real medical cases: medical document ingestion and extraction; specialist-trained medical AI agents; multi-agent verification and confidence scoring; source-level traceability; structured medical intelligence generation; medical and legal document synthesis; and case-specific chatbot over verified data. The next stage focuses on production and customer delivery.

Phase 1 - System reliability and monitoring (1–2 months): We will prepare the platform for production use by adding redundancy, failover, and real-time monitoring. This ensures high uptime, stable performance, and rapid failure detection. The goal is to meet enterprise-grade reliability standards and avoid disruptions during live case work.

Phase 2 - Secure cloud deployment and trial environment (1 month): We will deploy MedBrief in a secure, multi-environment cloud setup, separating development, trials, and production. A sandboxed 30-day free trial will allow customers to test the platform on real-world cases with strict data isolation, encryption, and automatic data deletion.

Phase 3 - End-user product experience (1–2 months): We will launch a role-based web application for lawyers, medical reviewers, and administrators. Users will be able to upload cases, view extracted data, review confidence scores, explore traceability links, and interact with the case through the chatbot. Training materials will enable users to become productive quickly.

Go-to-market assets: As we work through the above phases, we will also develop demo cases, security documentation, and ROI tools to support founder-led sales and customer trials.

Throughout the process, we will incorporate customer feedback and iterate. Early customer feedback will drive continuous improvements to accuracy, workflows, and pricing as we move toward product-market fit.

Describe the user experience for your product/service.

MedBrief is designed to integrate seamlessly into existing medical and legal workflows while removing the most time-consuming aspects of medical evidence review. The user experience prioritizes clarity, verification, and speed.

For physicians, the experience begins with document submission. Medical records are uploaded through a secure web interface. Once processing begins, MedBrief automatically extracts and structures medical data. The system generates the report, highlights extracted facts, and flags areas of uncertainty. Physicians review, confirm, or correct these sections, allowing them to focus on clinical judgment rather than manual compilation. Every verified data point remains linked to its original source in the medical record.

For legal professionals, the experience centers on working directly with verified medical evidence. Lawyers access structured medical summaries and timelines generated from the whole record, not just physician reports. Each fact is traceable to the underlying document, enabling confidence in both internal analysis and external advocacy. Lawyers can explore the case through a conversational interface (chatbot) that generates draft legal summaries and argument outlines grounded in the verified medical data. When new questions arise, users can query the record without re-reading hundreds of pages.

Throughout the experience, MedBrief emphasizes transparency. Confidence scores are visible. Uncertain data is clearly marked. Nothing is inferred silently. These features allow users to trust the system without surrendering professional control.

The overall experience is faster, calmer, and more deliberate. Tasks that previously took weeks, waiting for reports, reviewing records, and reconciling inconsistencies, are completed in minutes or hours. Users spend less time searching for information and more time making informed decisions, whether in clinical or legal contexts. MedBrief changes how efficiently and reliably professionals can work with medical evidence.

How have you validated that this product/service is possible to build, implement, and maintain?

We have validated the feasibility of MedBrief by building and testing the highest-risk components first and confirming that each core subsystem functions reliably in practice.

The most technically challenging aspect of the product is medical document extraction. We identified this early as the primary risk, since all downstream reporting and legal reasoning depend on the quality of the extracted data. To address this, we built a working prototype of MedBrief's agentic data ingestion and extraction pipeline and tested it on real-world medical records, including scanned documents, handwritten notes, and mixed-format files. This prototype has demonstrated that accurate, source-linked extraction is achievable at scale.

At the same time, we evaluated automated report generation using the verified extracted data and confirmed that the system can produce structured medical summaries. All outputs generated by the system are traceable to the exact page and location within the original medical record. When the system is uncertain, data is flagged rather than inferred. This design explicitly prevents hallucinations and supports legal and medical defensibility.

We also validated a conversational interface for lawyers that responds to user questions solely from verified report data and underlying source documents, rather than generating free-form answers.

Features yet to be implemented include security measures such as encrypted storage, authentication, two-factor authentication, and HIPAA-compliant data handling. These aspects can be implemented by leveraging existing, well-supported infrastructure rather than us building our own solutions. This reduces implementation risk and long-term maintenance.

Have potential users tested your product? Yes

Yes, potential users have tested core components of MedBrief, though the full end-to-end workflow has not yet been deployed.

Medical doctors have actively tested MedBrief's document extraction and verification capabilities. This was an intentional starting point. Given the principle of "garbage in, garbage out," we prioritized validating extraction accuracy before developing downstream features. Physicians and interns have reviewed extracted data and summaries generated from real medical records and provided feedback on accuracy, completeness, and clinical relevance. Their continued use of our service across multiple batches of records has confirmed that the extraction layer is practical and valuable.

Lawyers have tested MedBrief's conversational interface and legal argument generation tools in isolation. These tools were evaluated using verified medical data and case-specific inputs. Lawyers provided feedback on the usefulness of the generated summaries and argument structures, particularly for early case assessment and issue spotting. This testing helped validate the legal reasoning layer independently from the extraction pipeline.

However, the fully integrated workflow, in which raw medical records flow through extraction, physician verification, and legal analysis within a single user-facing platform, has not yet been tested end-to-end with legal users. We are currently refining the extraction and verification layers before running full pilot deployments.

This staged testing approach has allowed us to validate each critical component under real conditions while minimizing risk. Full workflow pilots with legal teams are planned as the next step, once the most technically sensitive components are reliable.

BUSINESS MODEL **(300 words each)**

Describe your planned revenue model

MedBrief operates as a B2B SaaS platform with usage-based pricing, designed for medical-legal work, accounting for high document volume, repeat usage, and high human review costs. Revenue is generated primarily through subscriptions and per-case fees paid by personal injury lawyers, insurance defense lawyers, and medical reviewers. MedBrief offers two pricing options depending on user needs: pay-per-case or annual subscription.

For individual practitioners, users can pay per case, typically billed at 100-200 pages of medical records per case, or choose an annual subscription. A standard annual individual plan supports approximately 20,000 pages, which corresponds to 100-150 cases per year for many practitioners. Additional usage is priced incrementally.

For small and mid-sized firms, MedBrief offers team plans with shared page quotas and volume-based discounts. Team subscriptions support multiple users drawing from a pooled allowance, with discounts of 15-25% depending on team size. This structure encourages internal adoption, supports collaboration across cases, and creates a natural path for expansion from individual usage to firm-wide deployment.

Larger organizations, including high-volume law firms and insurance carriers, are served through enterprise agreements. A typical enterprise contract includes a shared quota of approximately 1,000,000 pages per year and annual pricing around $300,000, with enterprise-grade support and security.

Pricing is anchored directly to the time and cost lawyers and physicians already incur in producing and reviewing medical reports. By replacing hours of manual review and external vendor spending, MedBrief delivers immediate ROI while allowing revenue to scale naturally with case volume and document complexity. Because automation drives marginal costs down as usage grows, the revenue model supports strong operating leverage as the platform scales.

Can you describe additional points of your business model in detail (i.e. strategy, cost/pricing structure, distribution channels, customer acquisition plan)? No/Yes

Customer Acquisition Plan

MedBrief uses a founder-led, trust-driven sales model that reflects how legal and medical professionals actually adopt new software. Customers are reached through direct outreach to personal injury lawyers, insurance defense lawyers, and medical reviewers, followed by a 30-day free trial that allows them to run up to 10 real cases through the system. This allows buyers to evaluate MedBrief in their own files rather than relying on demos, which is especially important in regulated environments.

This approach keeps customer acquisition costs low while producing high-quality conversions. Individual users typically convert within 2-4 weeks, with an estimated CAC of about $1,200. Small firm teams convert over 1-2 months with CAC around $3,500. Enterprise buyers require longer pilots and security reviews, with CAC modeled at approximately $50,000, but these contracts average $300,000 per year, making them highly attractive once the platform is embedded.

Because MedBrief becomes part of daily case workflows, retention is strong. We project that individual users generate approximately $10,500 per year and remain active for about four years, yielding a lifetime value of about $42,000. Team accounts average $40,000 per year over five years, while enterprise customers average $300,000 per year, with retention near 95%, resulting in over $2 million in lifetime value. This results in LTV-to-CAC ratios well above typical SaaS benchmarks.

Growth follows a land-and-expand model. Individual lawyers and medical reviewers start using MedBrief on their own cases; teams adopt it across firms; and organizations upgrade to enterprise contracts once MedBrief is embedded in their workflow. High switching costs, strong word-of-mouth, and immediate ROI support efficient, sustainable scaling.

Can you describe the type(s) of funding/resources and experts/strategic partners need to launch? No/Yes

To launch MedBrief, we require funding primarily for engineering, AI infrastructure, and security and compliance readiness. The platform's core technology is already validated, so capital will be used to make the system production-ready, reliable, and safe for use on real medical and legal data.

The largest use of funds will go toward engineering and AI infrastructure, including continued fine-tuning of specialist medical models, development of verification and confidence-scoring agents, and expansion of the secure document ingestion, chatbot, and report-generation pipelines. These components are what differentiate MedBrief from generic AI tools and must meet high reliability and accuracy standards.

A second major use of funds is cloud computing and storage. Medical records are large, sensitive, and used repeatedly across workflows. We will invest in encrypted storage, scalable processing infrastructure, and AI inference capacity to support customer trials and early commercial use.

A third priority is security and compliance. Funds will support HIPAA- and PHIPA-aligned data handling, access control, audit logging, and automated data deletion for trial environments. This is critical for earning trust from law firms, physicians, and insurers.

On the human resources side, the founding team already covers core expertise in law, medicine, and AI engineering. Early hires will focus on: an AI/ML engineer to improve verification and reasoning agents; a backend engineer to scale secure data pipelines; a full-stack engineer to build the end-user workflows and interface. We will also engage part-time compliance and privacy advisors to ensure the platform meets regulatory expectations as pilots begin.

Finally, we will continue to build an advisory network of personal injury lawyers, insurance defense attorneys, and specialist physicians to validate workflows, review outputs, and guide adoption. Strategic partnerships with legal service providers and insurance technology firms will become relevant as MedBrief scales.

How can the venture scale?

MedBrief scales through organizational expansion, contract growth, and market expansion.

First, MedBrief scales within organizations. Adoption typically begins with a small group of lawyers or medical reviewers who evaluate the platform on real cases. Once MedBrief proves its value by accelerating medical review and improving case quality, those users become internal champions who we can leverage to bring the platform to firm leadership, IT, and compliance for broader approval. From there, MedBrief is deployed across teams, offices, and practice groups, increasing contract size as it expands from pilot to firm-wide and enterprise contracts.

Second, MedBrief scales through volume-based economics. Law firms and insurers handle more cases as they grow, and each case involves large volumes of medical records. Because MedBrief is priced on pages and cases, revenue grows automatically as customers' workloads increase. This allows revenue to scale faster than headcount and keeps pricing aligned with customer value.

Third, MedBrief scales across adjacent regulated markets. The same core platform supports personal injury, insurance claims, disability, workers' compensation, and mass torts. These markets use similar medical data, require similar auditability, and face the same bottlenecks. Once MedBrief is trusted in one segment, it can be rolled out to others with minimal product change.

Finally, MedBrief scales through data and workflow lock-in. As firms rely on MedBrief for verified medical intelligence, switching becomes costly because historical cases, audit trails, and workflows are embedded in the system. This creates long-term retention and predictable revenue growth.

TEAM (300 words each)

Describe the composition of your venture's team. Please list each member's role (e.g., marketing lead, product developer, operations coordinator) and the corresponding responsibilities. Do not include any individual names or contact information.

MedBrief is built by a multidisciplinary founding team with expertise across law, medicine, and software engineering, enabling them to design, validate, and deploy an AI tool at the intersection of medicine and law.

Product Strategy Lead: Co-founder 1 is a JD candidate at Penn Carey Law (L'27) with an undergraduate degree in computer science and chemistry. She oversees day-to-day operations, product development, and cross-functional coordination. Her role includes translating legal and medical requirements into technical specifications, maintaining communication with both legal and medical users, and guiding product direction.

Medical Partnerships & Clinical Lead: Co-founder 2 is a practicing physician with decades of experience in both private practice and hospital settings. He leads outreach to the medical community and facilitates communication with practicing doctors who test the product and provide feedback. He contributes clinical expertise critical to evaluating source medical documents, coordinating medical interns to verify extracted data, and assessing the accuracy of generated reports.

Legal Partnerships & Industry Lead: Co-founder 3 is a practicing lawyer and professor who runs an active law firm. This role connects MedBrief with legal practitioners and provides insight into real-world litigation workflows. Responsibilities include user discovery with law firms, validation of legal summaries and argument structures, and guidance on market adoption.

Engineering and Systems Architecture Lead: Co-founder 4 is an engineering student with extensive experience building, deploying, and maintaining production-grade software systems. Responsibilities include developing backend infrastructure, refining AI pipelines, and translating novel concepts into reliable, operational systems.

Advisors and Supporting Members: The team is supported by a financial advisor and an insurance-industry advisor who provide insight into insurer workflows and claims processes. Additional support includes medical interns who assist with verifying extracted data and a senior developer with decades of software engineering experience.

Why are you the right founder/why is this the right founding team for this venture?
MedBrief addresses a problem that exists at the intersection of medicine, law, and technology. We are the right founding team for this venture because we are a highly motivated group with industry experience supporting each domain. We believe in focusing on execution, iteration, and learning through use and user feedback.

From the start, our team has focused on developing workable prototypes for the system's most critical components. Rather than attempting to build a fully featured platform from day one, we have isolated key technical and workflow components, tested them independently, and refined them through repeated feedback cycles. This approach has allowed us to identify what genuinely saves time, what builds user trust, and what is not helpful in practice.

Fast iteration is central to how the team operates. We actively seek feedback from practitioners, deploy improvements quickly, and reassess product direction based on how the system performs in live environments. This process has included removing or reshaping features that did not deliver clear value and focusing on those that demonstrably improve accuracy, speed, or usability. The team remains well-versed in ongoing AI developments. We continuously evaluate new models, tools, and techniques as benchmarks against which we compare our system, or as potential tools for our use.

The team’s interdisciplinary background enables this execution style. Legal insight informs how outputs must be structured and defended. Medical expertise underpins the system's real-world clinical reasoning. Engineering capability allows rapid prototyping, deployment, and revision without long development cycles. Ultimately, the team’s strength lies in its ability to connect, build, test, and iterate continuously.

We are proactive in engaging new users, quick to pursue opportunities for learning and growth, and disciplined in focusing effort to build and scale MedBrief.

What are the current gaps on your team and how are you addressing those gaps? Have you identified future team needs and what are they?

While MedBrief's founding team covers core expertise across law, medicine, and AI development, we recognize several capabilities that will become increasingly important as the venture transitions from development to deployment and scale.

The most immediate gap is dedicated sales and go-to-market execution. To date, customer engagement has been founder-led and driven by personal relationships, which has been effective for early validation and learning but is not scalable in the long term. As MedBrief enters pilot deployments and begins engaging larger law firms and medical practices, we will need team members with experience selling regulated software to professional services and enterprise buyers. In the near term, we are addressing this gap through direct founder involvement and advisor input as we refine product-market fit. In the long term, we plan to hire a sales or business development lead with experience in legal or healthcare technology.

A second gap focuses on regulatory and compliance. While the team can begin work on implementing secure data handling, security controls, and access management, we do not yet have a full-time compliance or privacy expert. As the platform scales and engages institutional customers, formal compliance oversight will become increasingly important. We plan to engage external compliance consultants initially and add in-house expertise as regulatory requirements expand.

As usage grows, additional engineering and infrastructure capacity will be required. The current engineering leadership is well-suited for development and early pilots, but scaling to higher document volumes and enterprise-grade reliability will require targeted backend and systems engineering support. We intend to address this through selective hiring once usage patterns and technical bottlenecks are better understood.

Importantly, none of these gaps undermines MedBrief's ability to build and validate the product today. They reflect the natural transition from a founding team to an operating company.

ACCOMPLISHMENTS & FUTURE PLANS (300 words each)

Describe the 3-5 key milestones that you have achieved to date that have moved the venture forward.

1. Validated core technical feasibility on real medical data: We built a working prototype of MedBrief's medical document extraction system and validated it on real-world records. To date, the system has processed over 600 pages of medical documentation spanning approximately 70 cases. These documents included a mix of formats and sources. Extraction outputs were reviewed and verified by practicing physicians, confirming that accurate, source-linked medical data can be extracted in practice.

2. Collaboration with practicing physicians outside of the founding team: Beyond extraction, MedBrief generated structured medical summaries and reports that were reviewed by physicians for use in their actual practice. Over the past three months, doctors have repeatedly submitted new batches of records for

processing and review. This ongoing use demonstrates that the outputs are not only technically correct but also clinically usable and time-saving in real reporting workflows.

3. Engagement with personal injury lawyers to validate legal relevance: We engaged practicing personal injury lawyers to review MedBrief's outputs and provide feedback on legal relevance, usability, and workflow fit. These conversations informed how medical data should be structured, summarized, and translated for litigation and claims contexts, shaping report formats and legal argument generation.

4. Reliable Document Extraction: We evaluated MedBrief's extraction and summarization performance against leading general-purpose and document-focused AI systems, including OpenAI-based models, Gemini (and Google's Med-focused AI tool), Qwen, and document AI platforms such as Landing AI. These comparisons highlighted MedBrief's advantage in accurate handling of long, unstructured medical records with source-level traceability.

What are your planned milestones over the next 6-12 months?

Over the next 6-12 months, MedBrief will focus on moving from a validated core technology to a production-ready platform with paying customers and repeatable deployment.

1. Production-ready platform (0-3 months)

We will complete system hardening, cloud deployment, and security architecture so MedBrief can be used on live customer data. This includes high-availability infrastructure, monitoring, encrypted storage, access controls, and a sandboxed 30-day trial environment. By the end of this phase, the platform will be safe and stable enough for external pilots.

2. End-user product launch (2-5 months)

We will release the first full web application for lawyers and medical reviewers. This includes case dashboards, document upload, verified medical summaries, traceability links, and the conversational interface. Training materials and onboarding flows will allow new users to become productive within hours.

3. Pilot programs with legal teams (4-8 months)

We will run structured pilots with personal injury and insurance defense firms using real cases. Success will be measured by time saved, reduction in manual review, and quality of medical and legal outputs. These pilots will be used to refine workflows, pricing, and packaging.

4. First paid deployments (6-9 months)

Following successful pilots, we will convert early users to paid individual and team plans. The goal is to establish initial recurring revenue and demonstrate willingness to pay across multiple customers.

5. Expansion and enterprise readiness (9-12 months)

We will expand usage within early firms and prepare for larger contracts by formalizing security documentation, compliance materials, and support processes. This positions MedBrief for insurance carriers and large firm deployments.

What risks have you identified? What steps have you taken to mitigate those risks?
Building MedBrief involves technical, regulatory, and go-to-market risks, especially because the platform handles sensitive medical and legal data. We address these risks directly through product design, security architecture, and our commercialization strategy.

MedBrief processes protected health information and attorney-client materials, so privacy and compliance are critical. To address this, customer data will be stored in fully isolated environments with end-to-end encryption, fine-grained access controls, and detailed audit logs. Trial accounts are sandboxed, customer data is never used for model training without consent, and trial data is automatically deleted at the end of the evaluation period. The system is designed from the outset to align with HIPAA, PHIPA, and legal confidentiality standards.

Incorrect or unsupported AI outputs could create legal or clinical risk. MedBrief avoids this by using a multi-agent verification architecture with field-level confidence scoring. Low-confidence or inconsistent data is flagged for human review, and the system is designed to never invent or infer facts when extraction is uncertain. All summaries, legal arguments, and chatbot responses are grounded in verified, traceable source data.

Lawyers and physicians are cautious about new technology. We address this through founder-led onboarding and a 30-day trial that allows users to test MedBrief on their own real cases, with strong data protection guarantees and immediate ROI.

Large AI providers and fast-moving startups may enter this space, but MedBrief already outperforms current solutions on what matters most: medical accuracy, verified extraction, and legal-grade traceability. These capabilities are built into our core architecture and are not easy to copy. With additional funding and support, MedBrief is positioned to maintain and extend its lead in this market.

Together, these measures keep risk manageable while allowing MedBrief to operate safely and credibly in regulated legal and medical markets.

USE OF FUNDS **(300 words each)**
Provide a budget for how you would use award funds, if you were to win an award through the Startup Challenge.
If awarded funding through the Startup Challenge, MedBrief would use the funds to move from a validated prototype to secure, real-world pilot deployments with legal and medical users. Our immediate goal is to support live customer trials while hardening the platform for reliability, security, and compliance.

With a $75,000 award, we would allocate the funds approximately as follows:

- $18,000 to cloud and AI compute to support document ingestion, extraction, verification, and chatbot usage across thousands of pages of real medical records.
- $12,000 to security and compliance, including HIPAA and PHIPA consulting, audit logging, sandboxed trial environments, and automated data deletion.

- $30,000 to fund engineering work, including system reliability, monitoring, user interfaces, and the trial provisioning system.
- $8,000 for customer onboarding and trial infrastructure, including account provisioning, training materials, and CRM tools.
- $7,000 to support legal, accounting, and operational setup, such as data processing agreements, privacy policies, and basic corporate compliance. This level of funding enables MedBrief to run secure pilots with 10-30 firms, convert early users to paid plans, and prepare for enterprise deployment.

With a $10,000 award, funds would be used to validate demand and technical defensibility at a smaller scale.

- $4,000 to support cloud and AI compute to run hundreds of pages of real cases.
- $2,000 toward security basics, including encryption, access controls, and sandboxed storage.
- $3,000 would fund engineering and hosting to deploy a stable demo environment, and
- $1,000 would support customer trials and outreach.

This would allow MedBrief to onboard 5-10 lawyers on real cases, demonstrate accuracy and traceability, and collect early evidence of willingness to pay. In both scenarios, funds are used to strengthen the product, build trust, and reach clear commercial milestones, rather than for marketing spend.

How will these funds directly support your venture development efforts? Why is that the best use of these funds.

The Startup Challenge funds would be used to convert MedBrief from a technically validated system into a product that real legal and medical professionals can safely use on live cases. The most critical barriers to adoption in this market are not awareness or marketing, but trust, security, and reliability. These funds directly address those barriers.

Investment in cloud infrastructure and AI compute enables us to process real medical records at scale during customer trials, allowing users to evaluate MedBrief on their own cases rather than in demos. Funding for security and compliance ensures that all customer data is handled in sandboxed, encrypted environments with audit logs and automatic deletion. These conditions must be met before law firms and physicians can upload sensitive files. Engineering support is used to harden the platform, improve verification and traceability, and deliver a stable end-user experience that supports daily workflows.

This is the best use of funds because it targets the highest-leverage bottlenecks in bringing MedBrief to market. Without secure infrastructure, compliance readiness, and a reliable product, even strong customer interest cannot turn into real adoption. By focusing the award on these foundational capabilities, MedBrief can move quickly into paid pilots, generate real revenue, and build the trust needed to scale in regulated legal and medical markets.

ADDITIONAL INFO

Is there anything else you'd like to share with the judges?

MedBrief is being built in close collaboration with practicing lawyers and physicians who deal with these problems every day. We are not starting from a blank slate. We are already processing real medical records, generating verified summaries, and incorporating feedback from professionals who rely on this work in high-stakes settings. What we are missing is not technical ambition, but the opportunity to turn this momentum into a stable, trusted product that can be used at scale.

The Startup Challenge would allow us to do exactly that. With access to Venture Lab's resources, mentorship, and community, we can move faster, avoid costly missteps, and build MedBrief to meet the standards of the legal and medical institutions it serves. We believe this problem is both economically significant and socially meaningful, and we are committed to building a solution that improves how medical evidence is used to reach fair outcomes for injured individuals.

We would be excited to build that future as part of the Venture Lab community.